Figure 0.1: Comparing methods on a $500 \times 500 \times 500$ synthetic dense tensor engineered such that the factor matrices have a few rows with concentrated leverage scores. Each method is run 10 times from random initializations (the same 10 initializations were used across all methods).